

4/24/2023

### Actor-Critic Algorithm

- Actor-critic Algorithm – separate V/Q value calculation and policy learning. The difference is that both get updated simultaneously.
- **Sidenote:** Expanding TD equation to infinity converges to Monte-Carlo equations
- Actor-critic algorithm utilizes the preference metric (from MAB problem – softmax/gradient policy)

$$\pi(a_t|s_t) = \frac{e^{H(s_t, a_t)}}{\sum_{a \in A} e^{H(s_t, a)}}$$

$$H(a_t) = H(a_t) + \alpha \frac{\partial E[R]}{\partial H_a}$$

$$H(a) = H(a) + \alpha [R_a - \bar{R}] [1'_{aa} - P_{a'}]$$

- In summary, two equations that are updated every step. Actor, which is the policy, and then the critic, which is the V or Q value.
  - o Critic:

$$V(s_t) = V(s_t) + \alpha \cdot \delta_t$$

- where  $\delta_t$  is the TD-0 error:  $\delta_t = [R_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$

- o Actor:

$$H(s_t, a_t) = H(s_t, a_t) + \beta \cdot \delta_t \cdot (1 - \pi(a_t|s_t))$$

- From Prof. Mahdi's HW3

#### Tabular Actor-Critic Algorithm

$V(s)=0, H(s,a)=0$ , for all  $s \in S, a \in A$ .

Repeat for  $N$  episodes

• Start from a random state  $s_0 \in S, t=0$

While  $t < T$  (episode length).

- Select action:  $a_t \sim \pi(\cdot | s_t)$  :  $\pi(a|s) = \frac{e^{H(s,a)}}{\sum_{a' \in A} e^{H(s,a' )}}$
- Take action  $a_t$ , move to state  $s_{t+1}$  and observe  $R_{t+1}$ .
- $\delta_t = R_{t+1} + \gamma V(s_{t+1}) - V(s_t)$
- $V(s_t) = V(s_t) + \alpha \delta_t$
- $H(s_t, a_t) = H(s_t, a_t) + \beta \delta_t (1 - \pi(a_t | s_t))$
- $t = t+1$

- Very fast convergence due to the gradient being used directly compared to other algorithms (policy + value iteration, SARSA + Q-learning, etc.)
- Also model-free, similar to SARSA + Q-learning
- Any policy can be used (doesn't have to be softmax). The only difference would have to be that the gradient would have to be computed accordingly, leading to a difference in the H or preference equation.

### Summary

- MDP -  $\langle S, A, R, P(s'|s, a) \rangle$
- If MDP is known, model-based policies
  - o Policy iteration
  - o Value iteration
- If MDP is not known, model-free policies
  - o Monte-Carlo
    - Generate a bunch of trajectories, estimate V values based on  $G_{ts}$
  - o Temporal Difference learning
    - Instead of going depth-wise, choose single action (or more, depending on TD length)
      - SARSA
      - Q-learning
  - o Actor-Critic
    - Similar to TD, but separation between policy and value function
- All these algorithms were tabularized, we had discrete tables. A real problem might not be tabularized. Actions, states, can be continuous (for example, car position/speed in Cartpole). Tabular algorithms no longer work ~ would converge too slowly/consume too much memory.
- **One solution:** approximate and discretize the action/state spaces. Tradeoff between closeness to true state space and time to converge.

### Continuous RL Algorithms

- **Option 1:** Instead of learning  $Q(s, a)$ , learn  $\theta$  (set of constants?) i.e.  $Q_\theta(s, a)$
- **Option 2:** DNN – states and action as input, Q value output. Another network can be used for the policy. Typically, 2-3 layers are enough to get good estimates. Question: How is the network trained?
- **Option 3:** Wavelets/Fourier transforms
- **Option 4:** Simple polynomials
  - o  $Q(s, a) = w_1 s \cdot a + w_2 s^2 \cdot a + w_3 s \cdot a^2$
  - o  $Q(s, a) = [w_1, w_2, w_3] \cdot [\phi_1(s, a), \phi_2(s, a), \phi_3(s, a)]^T$
  - o In this case, we're using known basis functions ( $\phi$ ) to express our Q-values as polynomials
  - o **Example:**
    - $V^\pi(s) \sim \tilde{V}(s, w) = w^T \cdot \phi(s)$
    - If  $\phi$  is chosen as indicator function i.e.  $1(s = s_1) = 1$  if  $s = s_1$ , 0 otherwise
      - $\tilde{V}(s, w) = w^T \cdot \phi(s) = [w_1, w_2, \dots, w_n] \cdot [1(s_1), 1(s_2), \dots, 1(s_n)]$  – this is the tabular case
    - Any class of  $\phi$  can be chosen

### Least Square Policy Iteration (LSPI)

- $\pi^0 \rightarrow PE \rightarrow V(s) \rightarrow PI \rightarrow \pi^1 \rightarrow \dots$
- Policy Evaluation PE:

$$V_{k+1}(s) = \sum_{s'} P(s'|s, \pi(s)) [R + \gamma V_k(s')]$$

Or

$$Q^\pi(s, a) = \sum_{s'} P(s'|s, a) [R + \gamma Q(s', \pi(s'))]$$

- Policy Improvement

$$\pi'(s) = \operatorname{argmax}_{a \in A} Q(s, a)$$

For model-free approach:

$$Q^\pi = R + \gamma \cdot M Q^\pi$$

Solving for Q,

$$Q^\pi = (I - \gamma M)^{-1} \cdot R$$

If Q is large, not possible to compute inverse in reasonable amount of time

So we apply the approximation principles as follows:

$$Q(s, a; w) = \sum_j \phi_j(s, a) \cdot w_j$$

See paper: <https://users.cs.duke.edu/~parr/jmlr03.pdf>

- Works well when state space is continuous,, but not as much when action nspace is continuous